

DSpectra_2nd_fraction

February 17, 2018

```
In [1]: # Import modues
# -*- encoding: utf-8 -*-
import numpy as np
import matplotlib as mpl
from matplotlib import rc
import math
import pandas as pd
import os
import itertools
import scipy
from scipy import stats
from scipy import ndimage
import seaborn as sns

import matplotlib as mpl
from matplotlib import cm
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

from statsmodels.stats.descriptivestats import sign_test
from statsmodels.stats.weightstats import zconfint
from statsmodels.stats.weightstats import *

from skimage import measure
from scipy import ndimage
from scipy import misc

from scipy.stats.stats import pearsonr, spearmanr
from collections import Counter

# from pandas import ExcelWriter
from sklearn.cross_validation import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error
from sklearn.ensemble import RandomForestRegressor
```

```

import statsmodels.stats.api as sm
from sklearn import cross_validation, datasets, linear_model, metrics
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import OneHotEncoder

from sklearn.cluster import KMeans
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, accuracy_score, f1_score, precision_s
from sklearn import cross_validation, datasets, grid_search, linear_model, metrics

from scipy.optimize import curve_fit
from scipy import signal

from sklearn import random_projection
from sklearn.decomposition import RandomizedPCA
from sklearn.decomposition import PCA
from sklearn import manifold
from sklearn.cluster import KMeans
import pickle
sns.set_style("whitegrid")
sns.set_palette('Accent')

rc('font', family='Arial') # change font for russian

% matplotlib inline
print 'Import Ready'

```

Import Ready

```

/usr/local/lib/python2.7/dist-packages/sklearn/cross_validation.py:41: DeprecationWarning: This
    "This module will be removed in 0.20.", DeprecationWarning)
/usr/local/lib/python2.7/dist-packages/statsmodels/compat/pandas.py:56: FutureWarning: The panda
    from pandas.core import datetools
/usr/local/lib/python2.7/dist-packages/sklearn/grid_search.py:42: DeprecationWarning: This modul
    DeprecationWarning)

```

```

In [2]: # Help function
def grain_to_num(x):
    if x == 'grain':
        return 5000
    elif x == 'rawgrain':
        return 7000
    else:
        return int(x)

```

```

markers = ['o', 's', 'd', 'v', 'h', '>', '<', '*', '^', '+']
markers_line = ['-o', '-s', '-d', '-v', '-h', '->', '-<', '-*']

cmap = plt.get_cmap('Accent')
colors = [cmap(i) for i in np.linspace(0, 1, 10)][1:]

In [38]: # Load datasets
spe_df_npk_prep = pd.read_pickle('spe_df_cl_type') # load smooth spectra

In [39]: # Drop useless data
spe_df_npk_prep = spe_df_npk_prep.loc[(spe_df_npk_prep.fraction != '100.dry') &
                                       (spe_df_npk_prep.fraction != '500.dry')]

In [40]: # Class is disbalanced
print Counter(spe_df_npk_prep.fraction)

Counter({'100': 29, '500': 26, 'grain': 24, 'rawgrain': 10})

In [41]: # Add data to rawgrain class with bootstrap (create 20 false data to rawgrain)
add_rawgrain_prep = spe_df_npk_prep[spe_df_npk_prep['fraction'] == 'rawgrain']

add_rawgrain_prep = add_rawgrain_prep.loc[:, [u'base_s', u'Si', u'P',
                                              u'S', u'Cl', u'K',
                                              u'Ca', u'Ti', u'Mn',
                                              u'Fe', u'Sr', u'Mo_Coh', u'Mo']]

mi = add_rawgrain_prep.min(axis=0)
ma = add_rawgrain_prep.max(axis=0)
all_d = []
for j in xrange(len(mi)):
    all_d.append(np.random.randint(mi[j], ma[j], 20))
c = add_rawgrain_prep.columns.tolist()
false_rawgrain_prep = pd.DataFrame(np.transpose(np.array(all_d)), columns=c)
false_rawgrain_prep['fraction'] = 'rawgrain'

spe_df_npk_prep = pd.concat([spe_df_npk_prep, false_rawgrain_prep])
print Counter(spe_df_npk_prep.fraction)

Counter({'rawgrain': 30, '100': 29, '500': 26, 'grain': 24})

In [42]: # Use normalization
fraction_labels = spe_df_npk_prep.fraction
ftype_name = Counter(fraction_labels).keys()
shifr = dict(zip(ftype_name, xrange(len(ftype_name))))
print 'Shifr fraction: ', shifr
fraction_labels = [shifr[x] for x in fraction_labels]
print 'Fraction: ', Counter(fraction_labels)

```

```

data_smooth = spe_df_npk_prep[[u'base_s', u'Si', u'P',
                                u'S', u'Cl', u'K',
                                u'Ca', u'Ti', u'Mn',
                                u'Fe', u'Sr', u'Mo_Coh', u'Mo']]

data_prep_r = (data_smooth - data_smooth.mean()) / (data_smooth.max() - data_smooth.min())

data_prep_z = (data_smooth - data_smooth.mean()) / data_smooth.std()

Shifr fraction: {'100': 0, 'rawgrain': 1, 'grain': 2, '500': 3}
Fraction: Counter({1: 30, 0: 29, 3: 26, 2: 24})

In [43]: # Check the shape
print data_smooth.shape, data_prep_r.shape, data_prep_z.shape

(109, 13) (109, 13) (109, 13)

In [48]: def classif_art3(all_data, all_labels, tdata=''):
    # We will evaluate different approaches to classification with the selection of parameters
    '''
    Function for optimize classification and calculate main quality metrics.
    Return:
        outputs with data
    Parameters:
        train_data
        train_labels
        test_data - data for test
        test_labels - labels for test
        tdata - type of data, if needed
    '''

    # Split to train and test
    print Counter(all_labels)
    train_data, test_data, train_labels, test_labels = cross_validation.train_test_split(
        all_data, all_labels, test_size=0.2, train_size=0.8, stratify=all_labels)

    print Counter(train_labels)
    print Counter(test_labels)
    # CV startegy and mertics
    # - , 20 %
    cv = cross_validation.StratifiedShuffleSplit(train_labels, n_iter = 10, test_size = 0.2)
    cv_metrics = cross_validation.StratifiedShuffleSplit(all_labels, n_iter = 10, test_size = 0.2)
    metrics_all = ['accuracy', 'precision_macro', 'f1_macro', 'recall_macro']
    metriks_names = ['accuracy', 'precision_macro', 'f1_macro', 'recall_macro']

    # 1 - Linear classification with gradient descent
    print '\t\tLinesr classif:'

```

```

log_regressor = linear_model.SGDClassifier(penalty='none', shuffle=True, random_state=0)
# grid search
parameters_grid = {
    'loss': ['hinge', 'log', 'squared_loss', 'modified_huber'],
    'fit_intercept': [True, False],
    'max_iter': np.linspace(1000, 15000, 5, dtype=int), #
}
grid_cv = grid_search.GridSearchCV(log_regressor, parameters_grid, scoring = 'f1_max')
grid_cv.fit(train_data, train_labels)
print 'Best with grid search: '
print '\ttestimator: ', grid_cv.best_estimator_
print '\tscore: ', grid_cv.best_score_
print '\tparameters: ', grid_cv.best_params_
print '\tfeatures importance: '
a = train_data.columns
imp = np.mean(np.abs(grid_cv.best_estimator_.coef_),0)
b = imp / np.sum(imp) * 100.
importances = pd.DataFrame(zip(a, b))
importances.columns = ['feature name', 'importance']
print '\t', importances.sort_values(by='importance', ascending=False)

print "\tDetailed classification report:"
y_true, y_pred = test_labels, grid_cv.best_estimator_.predict(test_data)
print '\t', classification_report(y_true, y_pred)

for i in xrange(len(metrics_all)):
    scor = metrics_all[i]
    scoring = cross_validation.cross_val_score(grid_cv.best_estimator_, all_data, all_labels,
                                                scoring=scoring, cv=cv_metricks)
    print 'Best SGDlin ' + metrics_names[i] + ' mean:{}, max:{}, min:{}, std:{}'.format(
        scor, scoring.max(), scoring.min(), scoring.std())

# 2 - RidgeClassifier
print
ridge_classifier = linear_model.SGDClassifier(penalty='l2', random_state=0)
print '\t\tRidge Classifier:'
# grid search
parameters_grid = {
    'loss': ['hinge', 'log', 'squared_loss', 'modified_huber'],
    'fit_intercept': [True, False],
    'max_iter': np.linspace(1000, 15000, 5, dtype=int),
    'alpha': np.linspace(0.0001, 1., num = 10) #
}
grid_cv = grid_search.GridSearchCV(ridge_classifier, parameters_grid, scoring = 'f1_max')
grid_cv.fit(train_data, train_labels)
print 'Best with grid search: '
print '\ttestimator: ', grid_cv.best_estimator_

```

```

print '\tscore: ', grid_cv.best_score_
print '\tparameters: ', grid_cv.best_params_
print '\tfeatures importamce: '
a = train_data.columns
imp = np.mean(np.abs(grid_cv.best_estimator_.coef_),0)
b = imp / np.sum(imp) * 100.
importances = pd.DataFrame(zip(a, b))
importances.columns = ['feature name', 'importance']
print '\t', importances.sort_values(by='importance', ascending=False)

print "\tDetailed classification report:"
y_true, y_pred = test_labels, grid_cv.best_estimator_.predict(test_data)
print '\t', classification_report(y_true, y_pred)

for i in xrange(len(metrics_all)):
    scor = metrics_all[i]
    scoring = cross_validation.cross_val_score(grid_cv.best_estimator_, all_data, a
                                                scoring = scor, cv = cv_metricks)
    print 'Best Ridge ' + metriks_names[i] + ' mean:{}, max:{}, min:{}, std:{}'.for

# 3 - Lasso (L1) linear regression
print
print '\t\tLasso Classifire:'
lasso_classifire = linear_model.SGDClassifier(penalty='l1', random_state=0)
parameters_grid = {
    'loss': ['hinge', 'log', 'squared_loss', 'modified_huber'],
    'fit_intercept': [True, False], # center data
    'max_iter': np.linspace(1000,15000, 5, dtype=int),
    'alpha': np.linspace(0.0001, 1., num = 10) #
}
grid_cv = grid_search.GridSearchCV(lasso_classifire, parameters_grid, scoring = 'f1')
grid_cv.fit(train_data, train_labels)
print 'Best with grid search: '
print '\testimator: ', grid_cv.best_estimator_
print '\tscore: ', grid_cv.best_score_
print '\tparameters: ', grid_cv.best_params_
print '\tfeatures importamce: '
a = train_data.columns
imp = np.mean(np.abs(grid_cv.best_estimator_.coef_),0)
b = imp / np.sum(imp) * 100.
importances = pd.DataFrame(zip(a, b))
importances.columns = ['feature name', 'importance']
print '\t', importances.sort_values(by='importance', ascending=False)

print "\tDetailed classification report:"
y_true, y_pred = test_labels, grid_cv.best_estimator_.predict(test_data)

```

```

print '\t', classification_report(y_true, y_pred)

for i in xrange(len(metrics_all)):
    scor = metrics_all[i]
    scoring = cross_validation.cross_val_score(grid_cv.best_estimator_, all_data, a
                                                scoring = scor, cv = cv_metricks)

    print 'Best L1 ' + metriks_names[i] + ' mean:{}, max:{}, min:{}, std:{}'.format
                                                scoring.min(), scoring.std())

# nonlinear Random Forest
print
rf_classifier = RandomForestClassifier(random_state=0)
print '\t\tRandom Forest:'
parameters_grid = {
    'n_estimators' : range(2, 100, 20),
    'max_features' : ['auto', 'sqrt', 'log2', None],
    'max_depth': [None] + range(2,13,5),
    'bootstrap' : [False, True],
    'class_weight': ['balanced', None]
}
grid_cv = grid_search.GridSearchCV(rf_classifier, parameters_grid, scoring = 'f1_ma
grid_cv.fit(train_data, train_labels)
print 'Best with grid search: '
print '\testimator: ', grid_cv.best_estimator_
print '\tscore: ', grid_cv.best_score_
print '\tparameters: ', grid_cv.best_params_
print '\tfeatures importamce: '
importances = pd.DataFrame(zip(train_data.columns, grid_cv.best_estimator_.feature_
importances.columns = ['feature name', 'importance']
print '\t', importances.sort_values(by='importance', ascending=False)
print "\tDetailed classification report:"
y_true, y_pred = test_labels, grid_cv.best_estimator_.predict(test_data)
print '\t', classification_report(y_true, y_pred)

for i in xrange(len(metrics_all)):
    scor = metrics_all[i]
    scoring = cross_validation.cross_val_score(grid_cv.best_estimator_, all_data, a
                                                scoring = scor, cv = cv_metricks)

    print 'Best RF ' + metriks_names[i] + ' mean:{}, max:{}, min:{}, std:{}'.format
                                                scoring.min(), scoring.std())

```

```

In [49]: # Calculate classification for fraction with data
#
list_df = [data_smooth, data_prep_r, data_prep_z]
list_names=['spe_df_npk_prep', 'data_prep_r', 'data_prep_z']
for i in xrange(len(list_df)):
    name = list_names[i]
    all_df_class = list_df[i]

```

```

print
print '\t Start new data classification ' + name
print
all_df_class.dropna(inplace=True)

all_df_class = all_df_class.astype(float, inplace=True)
classif_art3(all_data=all_df_class, all_labels=fraction_labels, tdata='fraction '+n
print

```

Start new data classification spe_df_npks_prep

```

Counter({1: 30, 0: 29, 3: 26, 2: 24})
Counter({1: 24, 0: 23, 3: 21, 2: 19})
Counter({0: 6, 1: 6, 2: 5, 3: 5})
      Linesr classif:

```

/usr/local/lib/python2.7/dist-packages/ipykernel_launcher.py:10: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>
Remove the CWD from sys.path while we load stuff.

Best with grid search:

```

      estimator:  SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='log', max_iter=15000, n_iter=None,
n_jobs=1, penalty='none', power_t=0.5, random_state=0, shuffle=True,
tol=None, verbose=0, warm_start=False)
      score:  0.707906436529
      parameters:  {'loss': 'log', 'max_iter': 15000, 'fit_intercept': True}
      features importance:
            feature name  importance
5              K    28.187664
2              P    20.938611
6             Ca    20.753671
4             Cl    11.098228
3              S     9.066928
10             Sr     2.968945
0          base_s     2.167351
9              Fe     1.272086
1              Si     1.262065
7              Ti     0.987836
12             Mo     0.669217
8              Mn     0.369684
11          Mo_Coh     0.257713

```


Detailed classification report:

	precision	recall	f1-score	support
0	0.44	0.67	0.53	6
1	1.00	1.00	1.00	6
2	0.71	1.00	0.83	5
3	0.00	0.00	0.00	5
avg / total	0.56	0.68	0.61	22

Best SGDlin accuracy mean:0.757575757576, max:0.969696969697, min:0.636363636364, std:0.10670798

Best SGDlin precision_macro mean:0.691620220058, max:0.975, min:0.516666666667, std:0.1750584610

Best SGDlin f1_macro mean:0.708349065905, max:0.970175438596, min:0.562834224599, std:0.13970372

Best SGDlin recall_macro mean:0.762698412698, max:0.96875, min:0.651785714286, std:0.10502345109

Ridge Classifire:

Best with grid search:

estimator: SGDClassifier(alpha=0.11120000000000001, average=False, class_weight=None, epsilon=0.1, eta0=0.0, fit_intercept=False, l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=15000, n_iter=None, n_jobs=1, penalty='l2', power_t=0.5, random_state=0, shuffle=True, tol=None, verbose=0, warm_start=False)

score: 0.573010887225

parameters: {'alpha': 0.11120000000000001, 'loss': 'log', 'max_iter': 15000, 'fit_intercept': True}

features importance:

	feature name	importance
5	K	25.535646
6	Ca	19.940428
4	Cl	19.938945
2	P	19.704375
10	Sr	4.384987
3	S	4.170666
7	Ti	1.689760
0	base_s	1.242397
9	Fe	1.209169
1	Si	0.948868
12	Mo	0.511172
8	Mn	0.483565
11	Mo_Coh	0.240022

Detailed classification report:

	precision	recall	f1-score	support
0	0.55	1.00	0.71	6
1	1.00	1.00	1.00	6
2	1.00	1.00	1.00	5
3	0.00	0.00	0.00	5
avg / total	0.65	0.77	0.69	22

Best Ridge accuracy mean:0.666666666667, max:0.757575757576, min:0.515151515152, std:0.081311562
 Best Ridge precision_macro mean:0.542567744004, max:0.632352941176, min:0.333333333333, std:0.08
 Best Ridge f1_macro mean:0.580577887926, max:0.673076923077, min:0.375, std:0.0973193668857
 Best Ridge recall_macro mean:0.672767857143, max:0.75, min:0.5, std:0.0829041201143

Lasso Classifier:

Best with grid search:

```

  estimator:  SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='modified_huber', max_iter=15000,
n_iter=None, n_jobs=1, penalty='l1', power_t=0.5, random_state=0,
shuffle=True, tol=None, verbose=0, warm_start=False)
  score:  0.717817270878
  parameters:  {'alpha': 0.0001, 'loss': 'modified_huber', 'max_iter': 15000, 'fit_interce
features importance:

```

	feature name	importance
5	K	28.182291
2	P	20.966581
6	Ca	20.746615
4	Cl	11.079199
3	S	9.069135
10	Sr	2.972935
0	base_s	2.165077
9	Fe	1.272165
1	Si	1.260527
7	Ti	0.988951
12	Mo	0.669804
8	Mn	0.369702
11	Mo_Coh	0.257017

Detailed classification report:

		precision	recall	f1-score	support
	0	0.44	0.67	0.53	6
	1	1.00	1.00	1.00	6
	2	0.71	1.00	0.83	5
	3	0.00	0.00	0.00	5
avg / total		0.56	0.68	0.61	22

Best L1 accuracy mean:0.730303030303, max:0.878787878788, min:0.636363636364, std:0.075999613357
 Best L1 precision_macro mean:0.678022763593, max:0.894444444444, min:0.515384615385, std:0.14100
 Best L1 f1_macro mean:0.674434217442, max:0.869543650794, min:0.562834224599, std:0.098688282165
 Best L1 recall_macro mean:0.734573412698, max:0.888888888889, min:0.630952380952, std:0.07706247

Random Forest:

Best with grid search:

```

  estimator:  RandomForestClassifier(bootstrap=False, class_weight=None, criterion='gini',

```

```

max_depth=2, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=42, n_jobs=1,
oob_score=False, random_state=0, verbose=0, warm_start=False)
score: 0.950300532801
parameters: {'max_features': 'auto', 'n_estimators': 42, 'bootstrap': False, 'max_depth':
features importance:
    feature name  importance
4             Cl   19.027399
5             K    16.991116
8            Mn   12.770117
6            Ca   12.321435
3             S   11.292295
9            Fe    9.280553
0         base_s    7.815014
2             P    7.778332
10           Sr    2.364107
7            Ti    0.359633
1            Si    0.000000
11         Mo_Coh    0.000000
12           Mo    0.000000
Detailed classification report:
              precision    recall  f1-score   support

0           1.00         0.67         0.80         6
1           1.00         1.00         1.00         6
2           0.71         1.00         0.83         5
3           1.00         1.00         1.00         5

avg / total         0.94         0.91         0.91        22

Best RF accuracy mean:0.890909090909, max:0.969696969697, min:0.787878787879, std:0.043281384415
Best RF precision_macro mean:0.897053571429, max:0.96875, min:0.790674603175, std:0.043341755288
Best RF f1_macro mean:0.886767533937, max:0.96862745098, min:0.78125, std:0.0445423958824
Best RF recall_macro mean:0.890873015873, max:0.972222222222, min:0.782738095238, std:0.04544311

Start new data classification data_prep_r

Counter({1: 30, 0: 29, 3: 26, 2: 24})
Counter({1: 24, 0: 23, 3: 21, 2: 19})
Counter({0: 6, 1: 6, 2: 5, 3: 5})
Linesr classif:
Best with grid search:
    estimator: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='log', max_iter=8000, n_iter=None,

```

```

n_jobs=1, penalty='none', power_t=0.5, random_state=0, shuffle=True,
tol=None, verbose=0, warm_start=False)
score: 0.95336996337
parameters: {'loss': 'log', 'max_iter': 8000, 'fit_intercept': True}
features importance:
    feature name  importance
4          Cl    24.193219
2           P    17.898992
5           K    13.975855
0       base_s     9.818958
3           S     7.837842
6          Ca     7.697182
9          Fe     6.522698
10         Sr     2.964303
8          Mn     2.616000
1          Si     2.390478
12         Mo     1.980904
7          Ti     1.493302
11       Mo_Coh     0.610266
Detailed classification report:
              precision    recall  f1-score   support

0           1.00         1.00         1.00         6
1           1.00         1.00         1.00         6
2           1.00         0.80         0.89         5
3           0.83         1.00         0.91         5

avg / total          0.96         0.95         0.95        22

Best SGDlin accuracy mean:0.930303030303, max:1.0, min:0.787878787879, std:0.0650633653139
Best SGDlin precision_macro mean:0.935109126984, max:1.0, min:0.805555555556, std:0.059994164289
Best SGDlin f1_macro mean:0.928986131902, max:1.0, min:0.794494720965, std:0.0632562679772
Best SGDlin recall_macro mean:0.927529761905, max:1.0, min:0.787202380952, std:0.0656220004953

Ridge Classifire:
Best with grid search:
    estimator: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='hinge', max_iter=4500, n_iter=None,
n_jobs=1, penalty='l2', power_t=0.5, random_state=0, shuffle=True,
tol=None, verbose=0, warm_start=False)
score: 0.970787545788
parameters: {'alpha': 0.0001, 'loss': 'hinge', 'max_iter': 4500, 'fit_intercept': True}
features importance:
    feature name  importance
4          Cl    30.799125
2           P    15.640904
5           K    14.732524

```

3	S	11.442922
0	base_s	9.579114
6	Ca	6.110882
9	Fe	4.892876
10	Sr	2.471119
7	Ti	1.361369
12	Mo	1.335452
8	Mn	0.718889
1	Si	0.623147
11	Mo_Coh	0.291678

Detailed classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	6
2	1.00	0.80	0.89	5
3	0.83	1.00	0.91	5
avg / total	0.96	0.95	0.95	22

Best Ridge accuracy mean:0.957575757576, max:1.0, min:0.848484848485, std:0.04535342287
 Best Ridge precision_macro mean:0.96152507215, max:1.0, min:0.86038961039, std:0.0419399296346
 Best Ridge f1_macro mean:0.956514861343, max:1.0, min:0.847435897436, std:0.0456097543457
 Best Ridge recall_macro mean:0.954712301587, max:1.0, min:0.842757936508, std:0.0470953605148

Lasso Classifire:

Best with grid search:

```

estimator: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=True, l1_ratio=0.15,
learning_rate='optimal', loss='log', max_iter=4500, n_iter=None,
n_jobs=1, penalty='l1', power_t=0.5, random_state=0, shuffle=True,
tol=None, verbose=0, warm_start=False)
score: 0.95336996337
parameters: {'alpha': 0.0001, 'loss': 'log', 'max_iter': 4500, 'fit_intercept': True}
features importance:

```

	feature name	importance
4	Cl	27.022628
2	P	19.875226
5	K	14.946657
0	base_s	9.369100
3	S	7.628318
6	Ca	7.172481
9	Fe	5.733864
10	Sr	2.791364
8	Mn	1.619879
1	Si	1.527793
12	Mo	1.045450
7	Ti	0.866189

```

11      Mo_Coh      0.401050
      Detailed classification report:
              precision    recall  f1-score   support

         0           1.00      1.00      1.00         6
         1           1.00      1.00      1.00         6
         2           1.00      0.80      0.89         5
         3           0.83      1.00      0.91         5

 avg / total           0.96      0.95      0.95        22

```

```

Best L1 accuracy mean:0.930303030303, max:1.0, min:0.787878787879, std:0.0636363636364
Best L1 precision_macro mean:0.934722222222, max:1.0, min:0.805555555556, std:0.0595855995425
Best L1 f1_macro mean:0.928212693309, max:1.0, min:0.794494720965, std:0.0627845193021
Best L1 recall_macro mean:0.926736111111, max:1.0, min:0.787202380952, std:0.0649441959432

```

```

      Random Forest:
Best with grid search:
      estimator: RandomForestClassifier(bootstrap=False, class_weight='balanced',
      criterion='gini', max_depth=None, max_features='auto',
      max_leaf_nodes=None, min_impurity_decrease=0.0,
      min_impurity_split=None, min_samples_leaf=1,
      min_samples_split=2, min_weight_fraction_leaf=0.0,
      n_estimators=42, n_jobs=1, oob_score=False, random_state=0,
      verbose=0, warm_start=False)
      score: 0.95211996337
      parameters: {'max_features': 'auto', 'n_estimators': 42, 'bootstrap': False, 'max_depth':
      features importance:
            feature name  importance
4            Cl      23.566913
5            K       14.752712
6            Ca      11.759159
2            P       10.341194
3            S        9.156011
0      base_s        8.704540
8            Mn        7.871371
9            Fe        6.780839
10           Sr         3.234122
7            Ti         1.361330
12           Mo         1.224135
1            Si         0.688572
11      Mo_Coh        0.559100

```

```

      Detailed classification report:
              precision    recall  f1-score   support

         0           1.00      1.00      1.00         6
         1           1.00      1.00      1.00         6
         2           1.00      0.80      0.89         5

```

	3	0.83	1.00	0.91	5
avg / total		0.96	0.95	0.95	22

Best RF accuracy mean:0.933333333333, max:1.0, min:0.818181818182, std:0.0465523984719
 Best RF precision_macro mean:0.937261904762, max:1.0, min:0.819444444444, std:0.0465252647644
 Best RF f1_macro mean:0.930867532236, max:1.0, min:0.813222724987, std:0.048209202012
 Best RF recall_macro mean:0.931696428571, max:1.0, min:0.810515873016, std:0.0485437711825

Start new data classification data_prep_z

Counter({1: 30, 0: 29, 3: 26, 2: 24})
 Counter({1: 24, 0: 23, 3: 21, 2: 19})
 Counter({0: 6, 1: 6, 2: 5, 3: 5})

Linesr classif:

Best with grid search:

estimator: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1, eta0=0.0, fit_intercept=False, l1_ratio=0.15, learning_rate='optimal', loss='squared_loss', max_iter=8000, n_iter=None, n_jobs=1, penalty='none', power_t=0.5, random_state=0, shuffle=True, tol=None, verbose=0, warm_start=False)

score: 0.941540824372

parameters: {'loss': 'squared_loss', 'max_iter': 8000, 'fit_intercept': False}

features importance:

	feature name	importance
3	S	24.579816
5	K	18.726601
4	Cl	18.470774
6	Ca	12.462850
2	P	8.865348
9	Fe	5.778613
0	base_s	4.378655
8	Mn	2.349314
10	Sr	1.642147
1	Si	1.633784
7	Ti	0.621826
12	Mo	0.271309
11	Mo_Coh	0.218964

Detailed classification report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	6
1	1.00	1.00	1.00	6
2	1.00	1.00	1.00	5
3	1.00	0.80	0.89	5
avg / total	0.96	0.95	0.95	22

Best SGDlin accuracy mean:0.890909090909, max:1.0, min:0.363636363636, std:0.178864903798
 Best SGDlin precision_macro mean:0.924605533356, max:1.0, min:0.632211538462, std:0.101892500714
 Best SGDlin f1_macro mean:0.892065701971, max:1.0, min:0.372474747475, std:0.176170682972
 Best SGDlin recall_macro mean:0.890079365079, max:1.0, min:0.37003968254, std:0.176496041605

Ridge Classifire:

Best with grid search:

```
estimator: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0, fit_intercept=False, l1_ratio=0.15,
learning_rate='optimal', loss='squared_loss', max_iter=8000,
n_iter=None, n_jobs=1, penalty='l2', power_t=0.5, random_state=0,
shuffle=True, tol=None, verbose=0, warm_start=False)
```

score: 0.941540824372

parameters: {'alpha': 0.0001, 'loss': 'squared_loss', 'max_iter': 8000, 'fit_intercept':

features importance:

	feature name	importance
3	S	24.318503
5	K	18.580339
4	Cl	18.405020
6	Ca	12.362111
2	P	9.218107
9	Fe	5.841083
0	base_s	4.410929
8	Mn	2.405032
10	Sr	1.709629
1	Si	1.665610
7	Ti	0.641261
12	Mo	0.252027
11	Mo_Coh	0.190350

Detailed classification report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	6
1	1.00	1.00	1.00	6
2	1.00	1.00	1.00	5
3	1.00	0.80	0.89	5
avg / total	0.96	0.95	0.95	22

Best Ridge accuracy mean:0.945454545455, max:1.0, min:0.848484848485, std:0.0402015126104
 Best Ridge precision_macro mean:0.953023989899, max:1.0, min:0.8625, std:0.0373364633136
 Best Ridge f1_macro mean:0.945767957318, max:1.0, min:0.852479757085, std:0.0388289406955
 Best Ridge recall_macro mean:0.943601190476, max:1.0, min:0.84623015873, std:0.0399137917093

Lasso Classifire:

Best with grid search:

```
estimator: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
```



```

eta0=0.0, fit_intercept=False, l1_ratio=0.15,
learning_rate='optimal', loss='squared_loss', max_iter=8000,
n_iter=None, n_jobs=1, penalty='l1', power_t=0.5, random_state=0,
shuffle=True, tol=None, verbose=0, warm_start=False)
score: 0.941540824372
parameters: {'alpha': 0.0001, 'loss': 'squared_loss', 'max_iter': 8000, 'fit_intercept':
features importance:
    feature name  importance
3             S   24.596737
5             K   18.645693
4            Cl   18.485669
6            Ca   12.421566
2             P    9.017770
9            Fe    5.832272
0        base_s    4.371286
8            Mn    2.386926
1            Si    1.675599
10           Sr    1.666248
7            Ti    0.588801
12           Mo    0.199693
11        Mo_Coh    0.111738
Detailed classification report:
              precision    recall  f1-score   support

0           0.86         1.00         0.92         6
1           1.00         1.00         1.00         6
2           1.00         1.00         1.00         5
3           1.00         0.80         0.89         5

avg / total           0.96         0.95         0.95        22

Best L1 accuracy mean:0.921212121212, max:1.0, min:0.69696969697, std:0.0848484848485
Best L1 precision_macro mean:0.936527777778, max:1.0, min:0.797222222222, std:0.0598506371347
Best L1 f1_macro mean:0.920741887878, max:1.0, min:0.691062801932, std:0.0859173742316
Best L1 recall_macro mean:0.918998015873, max:1.0, min:0.687003968254, std:0.0869531200393

Random Forest:
Best with grid search:
estimator: RandomForestClassifier(bootstrap=True, class_weight='balanced',
criterion='gini', max_depth=None, max_features='auto',
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=62, n_jobs=1, oob_score=False, random_state=0,
verbose=0, warm_start=False)
score: 0.943838734795
parameters: {'max_features': 'auto', 'n_estimators': 62, 'bootstrap': True, 'max_depth':
features importance:

```

	feature name	importance
4	Cl	19.941646
5	K	12.788595
6	Ca	12.696793
8	Mn	9.940975
2	P	9.382735
0	base_s	8.888559
3	S	8.119889
9	Fe	5.999332
1	Si	4.078361
7	Ti	3.086870
10	Sr	2.953553
12	Mo	1.300135
11	Mo_Coh	0.822559

Detailed classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	6
2	1.00	1.00	1.00	5
3	1.00	1.00	1.00	5
avg / total	1.00	1.00	1.00	22

Best RF accuracy mean:0.906060606061, max:1.0, min:0.787878787879, std:0.0566107324008
 Best RF precision_macro mean:0.908043831169, max:1.0, min:0.782738095238, std:0.059356867135
 Best RF f1_macro mean:0.901421613517, max:1.0, min:0.781512605042, std:0.0593924198908
 Best RF recall_macro mean:0.903521825397, max:1.0, min:0.782738095238, std:0.0598294278852

```
In [52]: # PCA
names = ['data_norm_r', 'data_norm_z']
datas = [data_prep_r, data_prep_z]
for ind in xrange(len(datas)):
    print 'Process data: ', names[ind]
    X_raw = datas[ind]
    labels = spe_df_npkp_prep.fraction

    print 'Labels counter: ', Counter(labels)
    x_data = X_raw.loc[:, X_raw.columns!='fraction']
    print 'Index len: ', len(x_data.index)
    feature_names = Counter(labels).keys()

    transformer = RandomizedPCA(n_components=2)
    X_2d = transformer.fit_transform(x_data)
    X_2d_norm = (X_2d - X_2d.mean()) / (X_2d.max() - X_2d.min())
    if ind == 1:
```

```

X_2d_norm = (X_2d - X_2d.mean()) / X_2d.std()
print 'After len: ', X_2d_norm.shape
# over 99.9% variance captured by 2d data
print transformer.explained_variance_ratio_

# do clustering
print 'n_cluster = ', len(Counter(labels))
estimator = KMeans(n_clusters=len(Counter(labels)), init='k-means++', n_init=10)
estimator.fit(X_2d_norm)

labels_t = estimator.labels_
print 'estimator labels: ', Counter(labels_t)

label_color = [colors[l] for l in labels_t]

title_font = {'fontname': 'Arial', 'size': '16', 'color': 'black', 'weight': 'normal',
              'verticalalignment': 'bottom'} # Bottom vertical alignment for more space

fig, ax = plt.subplots()
ax.scatter(X_2d_norm[:,0], X_2d_norm[:,1], c=label_color, s=50)
ax.scatter(estimator.cluster_centers_[:,0], estimator.cluster_centers_[:,1], marker=

# Calculate all mean for labels
shifr2 = {
    '100': u'pressed powder 100 mkm',
    '500': 'pressed powder 500 mkm',
    'rawgrain': 'granules',
    'grain': 'pressed granules'
}

all_means = []
for f in feature_names:
    x_m = X_2d_norm[labels==f,0].mean()
    y_m = X_2d_norm[labels==f,1].mean()
    all_means.append([round(x_m,3), round(y_m, 3), shifr2[f]])

print all_means

for k in Counter(labels_t).keys():
    x = X_2d_norm[:,0]
    y = X_2d_norm[:,1]
    text_now = ''

    plt.text(
        x[labels_t==k].mean(),
        y[labels_t==k].mean(),
        text_now,
        horizontalalignment='center',

```

```

        bbox=dict(alpha=.5, edgecolor='w', facecolor='w'),
        **title_font
    )

    plt.draw()
    plt.savefig(names[ind] + '_RandomizePCA_article3_2.png', dpi=300)
    plt.show()

```

Process data: data_norm_r

Labels counter: Counter({'rawgrain': 30, '100': 29, '500': 26, 'grain': 24})

Index len: 109

After len: (109, 2)


[0.90861113 0.02974156]

n_cluster = 4

estimator labels: Counter({3: 44, 1: 30, 0: 29, 2: 6})

[[-0.249, -0.029, u'pressed powder 100 mkm'], [0.573, -0.016, 'granules'], [-0.147, 0.057, 'pres

/usr/local/lib/python2.7/dist-packages/sklearn/utils/deprecation.py:58: DeprecationWarning: Class
warnings.warn(msg, category=DeprecationWarning)



Process data: data_norm_z

Labels counter: Counter({'rawgrain': 30, '100': 29, '500': 26, 'grain': 24})

Index len: 109

After len: (109, 2)

[0.86329941 0.05999724]

```

n_cluster = 4
estimator labels: Counter({3: 40, 1: 33, 0: 30, 2: 6})
[[-0.963, -0.113, u'pressed powder 100 mkm'], [2.147, -0.116, 'granules'], [-0.505, 0.278, 'pres

/usr/local/lib/python2.7/dist-packages/sklearn/utils/deprecation.py:58: DeprecationWarning: Clas
warnings.warn(msg, category=DeprecationWarning)

```

